

---

# Deep Neural Networks

Apprentissage par réseaux de neurones artificiels

ANNEN Rayane, MARTINS Alexis



12.04.2023

## Contents

<b>Digit recognition from raw data</b>	<b>2</b>
Shallow Neural Network . . . . .	3
Hyper-parameters . . . . .	3
<b>Digit recognition from features of the input data</b>	<b>3</b>
Shallow Neural Network . . . . .	3
Hyper-parameters . . . . .	3
<b>Convolutional neural network digit recognition</b>	<b>4</b>
Deep Convolutional Neural Network . . . . .	4
<b>Experiments</b>	<b>5</b>
Raw data . . . . .	5
100 neurones . . . . .	5
300 neurones . . . . .	6
600 neurones . . . . .	7
Features-based (HOG) . . . . .	8
PIX_P_CELL 4, orientation 8, 100 neurones . . . . .	8
PIX_P_CELL 4, orientation 8, 300 neurones . . . . .	9
PIX_P_CELL 4, orientation 4, 100 neurones . . . . .	10
PIX_P_CELL 7, orientation 8, 100 neurones . . . . .	11
CNN . . . . .	12
25 neurones . . . . .	12
250 neurones . . . . .	13
500 neurones . . . . .	14
<b>General questions</b>	<b>15</b>

## Digit recognition from raw data

**What is the learning algorithm being used to optimize the weights of the neural networks? What are the parameters (arguments) being used by that algorithm? What cost function is being used ? please, give the equation(s)**

The algorithm used to optimize the weight is RMSprop (Root Mean Square Propagation).

The parameters are the following:

```
1 tf.keras.optimizers.RMSprop(
2     learning_rate=0.001,
3     rho=0.9,
4     momentum=0.0,
5     epsilon=1e-07,
6     centered=False,
7     weight_decay=None,
8     clipnorm=None,
9     clipvalue=None,
10    global_clipnorm=None,
11    use_ema=False,
12    ema_momentum=0.99,
13    ema_overwrite_frequency=100,
14    jit_compile=True,
15    name="RMSprop",
16    **kwargs
17 )
```

The following equations are used:

$$E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta) \left( \frac{\partial C}{\partial w} \right)^2$$

$$w_t = w_{t-1} - \frac{\eta}{\sqrt{E[g^2]_t}} \frac{\partial C}{\partial w}$$

where  $\eta$  is the learning rate,  $w_t$  the new weight,  $\beta$  is the moving average parameter,  $E[g]$  is the moving average of squared gradients and  $\frac{\partial C}{\partial w}$  is the derivative of the cost function with respect to the weight.

The cost function is the categorical cross-entropy loss function:

$$\text{CE} = -\frac{1}{N} \sum_{k=0}^N \log \vec{p}_i[y_i]$$

where  $N$  is the number of samples,  $\vec{p}_i$  is the neural network output and  $y_i$  is the target class index.

## Shallow Neural Network

For this experiment a simple shallow neural network is used. We use raw data to classify the digits.

### Hyper-parameters

Changed made to the model from the original: we reduced the number of neurons in the hidden layer from 300 to 100.

- Epochs: 10
- Hidden layers:
  - 100 neurons, reLU
- Output activation function: softmax.
- Batch size: 128
- Weights in the hidden layer:  $784 * 100 + 100 = 78400 + 100 = 78500$
- Weights in the output layer =  $10 * 100 + 10 = 1010$
- Total weights:  $78500 + 1010 = 79510$

## Digit recognition from features of the input data

### Shallow Neural Network

In this experiment, we use the Histogram of gradients (HOG) features to classify the digits.

### Hyper-parameters

HOG: - orientation count: 8 - pixels per cell: 4

- Epochs: 10
- Hidden layers:
  - 100 neurons, reLU
- Output activation function: softmax.
- Batch size: 128

- Weights in the hidden layer:  $392 * 200 + 200 = 78400 + 200 = 78600$
- Weights in the output layer =  $10 * 200 + 10 = 2010$
- Total weights :  $78600 + 2010 = 80610$

## Convolutional neural network digit recognition

### Deep Convolutional Neural Network

- Epochs: 10
- Batch size: 128
- Hidden layers:
  - Convolutional 2D: 5x5
  - MaxPooling 2D: pool size: 2x2
  - Convolutional 2D: 5x5
  - MaxPooling 2D: pool size: 2x2
  - Convolutional 2D: 3x3
  - MaxPooling 2D: 2x2
  - Flatten layer
  - Dense (25 neurons), activation function: reLU

Output: 10 neurons, activation function: softmax

Model complexity :

1	Layer (type)	Output Shape	Param #
2	=====	=====	=====
3	l0 (InputLayer)	[(None, 28, 28, 1)]	0
4	l1 (Conv2D)	(None, 28, 28, 9)	234
5	l1_mp (MaxPooling2D)	(None, 14, 14, 9)	0
6	l2 (Conv2D)	(None, 14, 14, 9)	2034
7	l2_mp (MaxPooling2D)	(None, 7, 7, 9)	0
8	l3 (Conv2D)	(None, 7, 7, 16)	1312
9	l3_mp (MaxPooling2D)	(None, 3, 3, 16)	0
10			
11			
12			
13			
14			
15			
16			
17			


```
18 flat (Flatten)          (None, 144)          0
19
20 l4 (Dense)              (None, 25)          3625
21
22 l5 (Dense)              (None, 10)          260
23
24 =====
25 Total params: 7,465
26 Trainable params: 7,465
27 Non-trainable params: 0
28 -----
```

## Experiments

### Raw data

#### 100 neurones

Test score: 0.07702907919883728  
Test accuracy: 0.9757000207901001




```

313/313 [=====] - 1s 2ms/step
array([[ 972,    0,    1,    1,    0,    0,    2,    1,    3,    0],
       [    0, 1126,    2,    0,    0,    2,    2,    1,    2,    0],
       [    2,    1, 1016,    1,    3,    0,    2,    3,    4,    0],
       [    1,    0,    7, 994,    0,    2,    0,    1,    3,    2],
       [    2,    0,    2,    1, 971,    0,    2,    1,    0,    3],
       [    4,    0,    0, 11,    1, 863,    7,    0,    5,    1],
       [    3,    3,    1,    1,    3,    3, 943,    0,    1,    0],
       [    0,    4, 10,    4,    3,    0,    0, 998,    3,    6],
       [    3,    0,    2,    5,    2,    2,    1,    3, 954,    2],
       [    2,    2,    0,    7, 13,    4,    1,    3,    2, 975]])

```

### 300 neurones

Test score: 0.06406917423009872  
 Test accuracy: 0.9811999797821045




```

313/313 [=====] - 1s 2ms/step
array([[ 972,    0,    1,    1,    0,    0,    2,    1,    3,    0],
       [    0, 1126,    2,    0,    0,    2,    2,    1,    2,    0],
       [    2,    1, 1016,    1,    3,    0,    2,    3,    4,    0],
       [    1,    0,    7, 994,    0,    2,    0,    1,    3,    2],
       [    2,    0,    2,    1, 971,    0,    2,    1,    0,    3],
       [    4,    0,    0, 11,    1, 863,    7,    0,    5,    1],
       [    3,    3,    1,    1,    3,    3, 943,    0,    1,    0],
       [    0,    4, 10,    4,    3,    0,    0, 998,    3,    6],
       [    3,    0,    2,    5,    2,    2,    1,    3, 954,    2],
       [    2,    2,    0,    7, 13,    4,    1,    3,    2, 975]])

```

### 600 neurones

Test score: 0.060618676245212555  
 Test accuracy: 0.9815000295639038





```


313/313 [=====] - 1s 2ms/step
array([[ 970,    1,    2,    0,    1,    0,    3,    1,    2,    0],
       [   0, 1128,    2,    1,    0,    1,    2,    1,    0,    0],
       [   4,    2, 1005,    3,    3,    0,    2,    5,    8,    0],
       [   1,    0,    4,  994,    0,    5,    0,    1,    3,    2],
       [   0,    0,    1,    1,  967,    0,    6,    1,    0,    6],
       [   2,    0,    0,    5,    1,  871,    6,    0,    5,    2],
       [   3,    2,    1,    1,    3,    2,  945,    0,    1,    0],
       [   2,    5,    7,    4,    0,    0,    0, 1002,    4,    4],
       [   1,    2,    3,    4,    3,    2,    1,    2,  955,    1],
       [   3,    3,    0,    5,   10,    2,    1,    4,    3,  978]])

```

## Features-based (HOG)

PIX\_P\_CELL 4, orientation 8, 100 neurones

Test score: 0.06830427795648575  
 Test accuracy: 0.9771999716758728




```

313/313 [=====] - 1s 1ms/step
array([[ 971,    0,    1,    0,    0,    2,    3,    2,    0,    1],
       [   4, 1117,    1,    3,    1,    1,    4,    2,    2,    0],
       [   2,    1, 1015,    2,    2,    0,    2,    5,    3,    0],
       [   0,    1,    4, 986,    0,   10,    0,    4,    5,    0],
       [   3,    1,    3,    0, 955,    0,    1,    1,    3,   15],
       [   2,    1,    0,   10,    0, 875,    4,    0,    0,    0],
       [   4,    2,    1,    0,    3,    3, 945,    0,    0,    0],
       [   0,    2,    7,    3,    5,    0,    0, 995,    3,   13],
       [   7,    0,    6,    8,    2,    2,    4,    4, 927,   14],
       [   1,    2,    1,    5,    7,    1,    0,    5,    1, 986]])

```

### PIX\_P\_CELL 4, orientation 8, 300 neurones

Test score: 0.09305289387702942  
 Test accuracy: 0.9696000218391418




```

313/313 [=====] - 0s 1ms/step
array([[ 966,    1,    2,    0,    0,    3,    4,    1,    3,    0],
       [   0, 1123,    3,    1,    0,    0,    4,    0,    4,    0],
       [   4,    7,  998,    4,    1,    0,    2,    8,    7,    1],
       [   0,    1,    2,  976,    0,   12,    0,    2,   16,    1],
       [   3,    2,    3,    0,  956,    0,    4,    2,    4,    8],
       [   4,    1,    0,   13,    0,  863,    5,    0,    5,    1],
       [   5,    1,    1,    0,    5,    4,  939,    0,    3,    0],
       [   2,    5,    6,    3,    7,    0,    0,  993,    3,    9],
       [   7,    2,    2,   11,    3,    3,    3,    3,  936,    4],
       [   5,    6,    0,   10,   18,    6,    1,   11,    6,  946]])

```

### PIX\_P\_CELL 4, orientation 4, 100 neurones

Test score: 0.06113172695040703  
 Test accuracy: 0.9801999926567078




```

313/313 [=====] - 0s 1ms/step
array([[ 972,    0,    1,    0,    0,    1,    3,    1,    1,    1],
       [   4, 1122,    1,    2,    0,    1,    2,    0,    3,    0],
       [   2,    4, 1017,    1,    1,    0,    2,    2,    3,    0],
       [   0,    1,    2,  987,    0,    6,    1,    3,   10,    0],
       [   1,    1,    1,    0,  962,    0,    0,    1,    3,   13],
       [   2,    1,    0,   12,    0,  869,    4,    0,    1,    3],
       [   4,    2,    1,    0,    3,    2,  945,    0,    1,    0],
       [   0,    4,    8,    4,    4,    0,    0,  988,    3,   17],
       [   6,    0,    4,    2,    1,    1,    1,    1,  949,    9],
       [   0,    3,    1,    3,    5,    1,    0,    4,    1,  991]])

```

### PIX\_P\_CELL 7, orientation 8, 100 neurones

Test score: 0.13366223871707916  
 Test accuracy: 0.9563999772071838



```


313/313 [=====] - 1s 2ms/step
array([[ 950,    1,    4,    1,    0,    1,   10,    1,    9,    3],
       [   0, 1109,    3,    2,    3,    1,    7,    1,    9,    0],
       [   3,    2, 989,   10,    6,    1,    2,    7,   11,    1],
       [   1,    1,    7, 940,    2,   25,    1,    7,   24,    2],
       [   0,    2,    2,    0, 953,    1,    1,    4,    5,   14],
       [   1,    0,    1,    8,    0, 865,    1,    0,   15,    1],
       [   4,    2,    0,    1,    8,   10, 927,    0,    4,    2],
       [   0,    5,   16,    5,    8,    0,    0, 957,   10,   27],
       [   2,    3,    2,    8,    4,   15,    3,    5, 924,    8],
       [   0,    3,    2,   11,   15,    5,    0,   12,   11, 950]])

```

## CNN

### 25 neurones

Test score: 0.03834224492311478  
 Test accuracy: 0.9879999756813049




```

313/313 [=====] - 1s 2ms/step
pred.shape = (10000, 10)
array([[ 975,    0,    0,    0,    0,    0,    1,    1,    3,    0],
       [   1, 1127,    2,    0,    0,    2,    0,    1,    2,    0],
       [   3,    2, 1017,    2,    0,    0,    0,    7,    1,    0],
       [   0,    0,    1, 1003,    0,    3,    0,    2,    1,    0],
       [   0,    0,    2,    0,  971,    0,    4,    2,    0,    3],
       [   0,    1,    0,    6,    0,  882,    1,    0,    1,    1],
       [   5,    2,    0,    0,    2,    6,  939,    0,    4,    0],
       [   0,    2,    2,    2,    0,    1,    0, 1020,    1,    0],
       [   4,    0,    2,    0,    0,    1,    0,    2,  964,    1],
       [   2,    3,    0,    3,    6,    3,    0,    6,    4,  982]])

```

### 250 neurones

Test score: 0.037620659917593  
 Test accuracy: 0.9884999990463257




```

313/313 [=====] - 1s 2ms/step
pred.shape = (10000, 10)
array([[ 975,    0,    1,    0,    0,    0,    1,    0,    2,    1],
       [   1, 1131,    2,    0,    0,    0,    1,    0,    0,    0],
       [   2,    0, 1028,    0,    0,    0,    0,    1,    1,    0],
       [   0,    0,    4,  999,    0,    3,    0,    2,    1,    1],
       [   0,    1,    3,    0,  964,    0,    4,    0,    3,    7],
       [   2,    0,    0,    5,    0,  882,    1,    0,    0,    2],
       [   7,    1,    1,    0,    1,    3,  945,    0,    0,    0],
       [   0,    1,   12,    0,    0,    0,    0, 1012,    1,    2],
       [   4,    0,    4,    2,    0,    2,    1,    2,  955,    4],
       [   0,    1,    0,    0,    5,    4,    0,    1,    4,  994]])

```

### 500 neurones

Test score: 0.04010535404086113  
 Test accuracy: 0.9871000051498413



```

313/313 [=====] - 1s 2ms/step
pred.shape = (10000, 10)
array([[ 973,    0,    0,    0,    0,    2,    2,    0,    2,    1],
       [   0, 1131,    1,    1,    0,    0,    1,    1,    0,    0],
       [   1,    3, 1008,   12,    1,    0,    0,    5,    1,    1],
       [   0,    0,    0, 1006,    0,    3,    0,    0,    1,    0],
       [   0,    1,    0,    0, 975,    0,    0,    0,    1,    5],
       [   1,    0,    0,   10,    0, 877,    1,    0,    0,    3],
       [   1,    2,    0,    1,    2,   13, 938,    0,    1,    0],
       [   0,    1,    5,    3,    1,    0,    0, 1010,    1,    7],
       [   0,    0,    0,   15,    1,    1,    0,    1, 953,    3],
       [   0,    0,    0,    2,    3,    2,    0,    1,    1, 1000]])

```

## General questions

**Do the deep neural networks have much more “capacity” (i.e., do they have more weights?) than the shallow ones? explain with one example**

Yes, deep neural networks generally have much more capacity than shallow ones, as they have more layers and consequently more weights.

Counterintuitively, it is observed that the shallow model has many more parameters than the deep model. This is because a shallow model is heavily interconnected, which increases the number of parameters.

We can use the example of the laboratory where the shallow model has 10 times more parameters than the deep one.